

Data Visualization

```
import numpy as np
import pandas as pd
import pandas as pd
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt

from google.colab import drive
import os
import pandas as pd

# Mount Google Drive
drive.mount('/content/drive', force_remount=True)

# Specify the folder path in your Google Drive
folder_path = '/content/drive/MyDrive/CIC-IDS- 2017' # Replace 'CIC-IDS-2017' with the actual folder name

# Get the list of file names in the folder
file_names = os.listdir(folder_path)

# Initialize an empty list to store DataFrames
dfs = []

# Loop through each file in the folder
for file_name in file_names:
    if file_name.endswith('.csv'):
        file_path = os.path.join(folder_path, file_name)
        # Read the CSV file into a DataFrame
        df = pd.read_csv(file_path)
        # Append the DataFrame to the list
        dfs.append(df)

# Concatenate the list of DataFrames into a single DataFrame
combined_df = pd.concat(dfs, ignore_index=True)

# Display the combined DataFrame
print(combined_df)
```

Mounted at /content/drive

| | Destination | Port | Flow Duration | Total Fwd Packets | \ |
|---|-------------|-------|---------------|-------------------|---|
| 0 | | 22 | 1266342 | 41 | |
| 1 | | 22 | 1319353 | 41 | |
| 2 | | 22 | 160 | 1 | |
| 3 | | 22 | 1303488 | 41 | |
| 4 | | 35396 | 77 | 1 | |

| | | | |
|-----------------------------|------------|-----------------------------|--------|
| ... | ... | ... | ... |
| 2830738 | 53 | 32215 | 4 |
| 2830739 | 53 | 324 | 2 |
| 2830740 | 58030 | 82 | 2 |
| 2830741 | 53 | 1048635 | 6 |
| 2830742 | 53 | 94939 | 4 |
| | | | |
| Total Backward Packets | | Total Length of Fwd Packets | \ |
| 0 | 44 | 2664 | |
| 1 | 44 | 2664 | |
| 2 | 1 | 0 | |
| 3 | 42 | 2728 | |
| 4 | 2 | 0 | |
| | | | |
| 2830738 | 2 | 112 | |
| 2830739 | 2 | 84 | |
| 2830740 | 1 | 31 | |
| 2830741 | 2 | 192 | |
| 2830742 | 2 | 188 | |
| | | | |
| Total Length of Bwd Packets | | Fwd Packet Length Max | \ |
| 0 | 6954 | 456 | |
| 1 | 6954 | 456 | |
| 2 | 0 | 0 | |
| 3 | 6634 | 456 | |
| 4 | 0 | 0 | |
| | | | |
| 2830738 | 152 | 28 | |
| 2830739 | 362 | 42 | |
| 2830740 | 6 | 31 | |
| 2830741 | 256 | 32 | |
| 2830742 | 226 | 47 | |
| | | | |
| Fwd Packet Length Min | | Fwd Packet Length Mean | \ |
| 0 | 0 | 64.975610 | |
| 1 | 0 | 64.975610 | |
| 2 | 0 | 0.000000 | |
| 3 | 0 | 66.536585 | |
| 4 | 0 | 0.000000 | |
| | | | |
| 2830738 | 28 | 28.000000 | |
| 2830739 | 42 | 42.000000 | |
| 2830740 | 0 | 15.500000 | |
| 2830741 | 32 | 32.000000 | |
| 2830742 | 47 | 47.000000 | |
| | | | |
| Fwd Packet Length Std | | min_seg_size_forward | Active |
| Mean \ | 109.864573 | ... | 32 |
| 0 | | | |
| 0.0 | | | |

| 1 0. 0 | | 109.864573 | ... | | 32 |
|--------------------|------------|------------|------------|-----------|------|
| 2 0. 0 | | 0.000000 | ... | | 32 |
| 3 0. 0 | | 110.129945 | ... | | 32 |
| 4 0. 0 | | 0.000000 | ... | | 32 |
| ... | | ... | ... | | ... |
| ... | | ... | ... | | ... |
| 2830738 0. 0 | | 0.000000 | ... | | 20 |
| 2830739 0. 0 | | 0.000000 | ... | | 20 |
| 2830740 0. 0 | | 21.920310 | ... | | 32 |
| 2830741 0. 0 | | 0.000000 | ... | | 20 |
| 2830742 0. 0 | | 0.000000 | ... | | 20 |
| Std \ Std | Active Std | Active Max | Active Min | Idle Mean | Idle |
| 0 | 0.0 | 0 | 0 | 0.0 | 0.0 |
| 1 | 0.0 | 0 | 0 | 0.0 | 0.0 |
| 2 | 0.0 | 0 | 0 | 0.0 | 0.0 |
| 3 | 0.0 | 0 | 0 | 0.0 | 0.0 |
| 4 | 0.0 | 0 | 0 | 0.0 | 0.0 |
| ... | ... | ... | ... | ... | ... |
| 2830738 | 0.0 | 0 | 0 | 0.0 | 0.0 |
| 2830739 | 0.0 | 0 | 0 | 0.0 | 0.0 |
| 2830740 | 0.0 | 0 | 0 | 0.0 | 0.0 |
| 2830741 | 0.0 | 0 | 0 | 0.0 | 0.0 |
| 2830742 | 0.0 | 0 | 0 | 0.0 | 0.0 |
| | Idle Max | Idle Min | Label | | |
| 0 | 0 | 0 | BENIGN | | |
| 1 | 0 | 0 | BENIGN | | |
| 2 | 0 | 0 | BENIGN | | |

```

3          0          0  BENIGN
4          0          0  BENIGN
...
2830738      0          0  BENIGN
2830739      0          0  BENIGN
2830740      0          0  BENIGN
2830741      0          0  BENIGN
2830742      0          0  BENIGN

[2830743 rows x 79 columns]

df.columns

Index([' Destination Port', ' Flow Duration', ' Total Fwd Packets',
       ' Total Backward Packets', 'Total Length of Fwd Packets',
       ' Total Length of Bwd Packets', ' Fwd Packet Length Max',
       ' Fwd Packet Length Min', ' Fwd Packet Length Mean',
       ' Fwd Packet Length Std', 'Bwd Packet Length Max',
       'Bwd Packet Length Min', ' Bwd Packet Length Mean',
       'Bwd Packet Length Std', 'Flow Bytes/s', ' Flow Packets/s',
       ' Flow IAT Mean', ' Flow IAT Std', ' Flow IAT Max', ' Flow IAT
Min',
       'Fwd IAT Total', ' Fwd IAT Mean', ' Fwd IAT Std', ' Fwd IAT
Max',
       ' Fwd IAT Min', 'Bwd IAT Total', ' Bwd IAT Mean', ' Bwd IAT
Std',
       ' Bwd IAT Max', ' Bwd IAT Min', 'Fwd PSH Flags', ' Bwd PSH
Flags',
       ' Fwd URG Flags', ' Bwd URG Flags', ' Fwd Header Length',
       'Bwd Header Length', 'Fwd Packets/s', ' Bwd Packets/s',
       ' Min Packet Length', ' Max Packet Length', ' Packet Length
Mean',
       ' Packet Length Std', ' Packet Length Variance', 'FIN Flag
Count',
       ' SYN Flag Count', ' RST Flag Count', ' PSH Flag Count',
       'ACK Flag Count', ' URG Flag Count', ' CWE Flag Count',
       'ECE Flag Count', ' Down/Up Ratio', ' Average Packet Size',
       'Avg Fwd Segment Size', ' Avg Bwd Segment Size',
       ' Fwd Header Length.1', 'Fwd Avg Bytes/Bulk', ' Fwd Avg
Packets/Bulk',
       ' Fwd Avg Bulk Rate', ' Bwd Avg Bytes/Bulk', ' Bwd Avg
Packets/Bulk',
       'Bwd Avg Bulk Rate', 'Subflow Fwd Packets', ' Subflow Fwd
Bytes',
       ' Subflow Bwd Packets', ' Subflow Bwd Bytes',
'Init_Win_bytes_forward',
       ' Init_Win_bytes_backward', ' act_data_pkt_fwd',
       ' min_seg_size_forward', 'Active Mean', ' Active Std', ' Active
Max',
       ' Active Min', 'Idle Mean', ' Idle Std', ' Idle Max', ' Idle

```

```

Min',
    'Label'],
    dtype='object')

set(list(df.dtypes))

{dtype('int64'), dtype('float64'), dtype('O')}

import pandas as pd

# Assuming df is the combined DataFrame
categorical_columns = [col for col in df.columns if df[col].dtype == "O"]

# Detect the label column dynamically
label_column = next((col for col in df.columns if 'label' in col.lower()), None)

if label_column:
    # Map "BENIGN" to itself and all other values to "INTRUSION"
    df[label_column] = df[label_column].map(lambda x: "BENIGN" if x == "BENIGN" else "INTRUSION")

    # Print the unique values in the modified label column
    print(df[label_column].unique())
else:
    print("Label column not found in the DataFrame.")

# Rest of your code using the categorical_columns variable...

['BENIGN' 'INTRUSION']

df_Categorical=df[categorical_columns]
df_Categorical

{"type":"dataframe","variable_name":"df_Categorical"}

df_Numerical=df.copy()
df_Numerical.drop(["Label"],axis=1,inplace=True)


df_Numerical

{"type":"dataframe","variable_name":"df_Numerical"}

df_Categorical["
Label"].value_counts(),sns.countplot(x=df_Categorical['Label'])

( Label
BENIGN      440031
INTRUSION    252672
Name: count, dtype: int64,
<Axes: xlabel=' Label', ylabel='count'>

```



```
for i in df_Numerical.columns:  
    #if df[i].unique()<  
    print(f" {i} Column has {df[i].nunique()} unique values and those  
are {df[i].unique()}\n\n")  
  
for i in df_Numerical.columns:  
    if df[i].isnull().sum()>0:  
        print(i)  
  
Flow Bytes/s  
  
df_Numerical["Flow Bytes/s"].isnull().sum()  
1008  
  
for index, row in df_Numerical.iterrows():  
    if pd.isna(row["Flow Bytes/s"]):  
        print(row["Flow Bytes/s"],df_Categorical["  
Label"].iloc[index])  
  
df_Numerical.dropna(inplace=True)
```

Infinity Valued Column Detection

```
np.all(np.isfinite(df_Numerical.values))
False

np.isinf(df_Numerical).values.sum()
578

infinite_column_count = {}
infinite_Value_index_num = []

for column in df_Numerical.columns:
    count = 0
    for index, value in enumerate(df_Numerical[column]):
        if not np.isfinite(value):
            count += 1
            infinite_Value_index_num.append(index)
    if count > 0:
        infinite_column_count[column] = count

infinite_Value_index_num = set(infinite_Value_index_num)

print("Counts of non-numeric (infinite) values in each column:")
print(f"{infinite_column_count}\n")

print("Index numbers of non-numeric (infinite) values:")
print(infinite_Value_index_num)

Counts of non-numeric (infinite) values in each column:
{'Flow Bytes/s': 289, ' Flow Packets/s': 289}

Index numbers of non-numeric (infinite) values:
{659968, 578052, 589320, 639498, 355852, 436748, 537101, 477199,
 637968, 21521, 672784, 569875, 3097, 657945, 21021, 9758, 583197,
 420896, 682527, 607267, 408613, 424997, 14375, 499749, 420908, 436780,
 352302, 548909, 481328, 666668, 572473, 389690, 441914, 608255,
 625657, 389694, 412738, 311364, 610886, 438343, 415816, 394315,
 568395, 523855, 191568, 349776, 660562, 636501, 63062, 24152, 361560,
 383577, 493145, 423004, 500825, 538715, 647258, 685662, 7780, 25702,
 147558, 335976, 366694, 393832, 431210, 555115, 26221, 48237, 13423,
 8816, 423023, 334962, 314994, 486003, 383606, 599670, 594040, 324217,
 416377, 425594, 29825, 348290, 681601, 26244, 51334, 340615, 455303,
 431241, 339082, 514184, 336015, 541842, 450707, 20632, 429209, 621721,
 625818, 407708, 466079, 639142, 393383, 426152, 649895, 683690, 15019,
 40620, 474797, 476340, 357557, 45238, 514230, 496312, 468666, 44731,
 345277, 540862, 468671, 528063, 581311, 648895, 399555, 224454,
 504521, 3789, 4301, 473805, 373968, 431825, 473809, 646354, 681166,
 684237, 608470, 415447, 612567, 382684, 28381, 143070, 388317, 644319,
 323809, 478946, 66787, 66788, 29413, 438503, 483048, 405737, 483052,
```

```
295661, 382702, 49903, 348912, 649453, 685806, 520435, 340725, 415989,
340215, 6904, 357112, 418041, 429815, 486140, 351485, 604917, 650487,
668416, 677625, 380675, 356616, 532749, 477457, 441106, 670483,
340758, 570648, 129306, 651547, 484125, 30495, 671523, 431397, 583467,
337709, 451373, 480047, 575794, 426293, 59190, 471355, 593726, 651583,
343873, 402754, 439105, 509249, 44358, 566599, 353610, 4939, 360266,
23375, 418131, 662868, 50517, 682838, 433496, 2395, 384866, 337252,
562534, 512874, 578924, 519533, 1392, 340849, 51569, 57202, 64368,
666997, 522617, 597370, 44411, 648573, 8064, 465794, 466820, 603524,
24456, 361352, 531336, 623497, 686986, 50577, 416146, 530326, 6551,
406935, 418712, 514970, 45472, 463780, 550313, 683946, 362414, 607665,
671666, 431539, 572339, 671161, 568250, 466363, 503231, 324034,
406979, 59844, 418761, 421324, 501197, 636365, 56271, 336336, 394704,
474577, 339413, 349142, 428503, 440792, 622039, 577498, 15835, 346587,
536028, 557049, 555999, 587231, 631770, 593890, 647656, 539117,
672750, 530927, 469489, 498673, 442867, 529908, 2041, 20991}
```

```
Benign=0
Intrusion=0
for i in infinite_Value_index_num:
    if df_Categorical["Label"].iloc[i]=="BENIGN":
        Benign+=1
    else:
        Intrusion+=1

print(f" Normal counts are {Benign}\n\n")
print(f" Intrusion counts are {Intrusion}")
```

```
Normal counts are 276
```

```
Intrusion counts are 13
```

```
df_Numerical.replace([np.inf, -np.inf], np.nan, inplace=True)
df_Numerical.dropna(inplace=True)

np.all(np.isinf(df_Numerical)),np.any(np.isnan(df_Numerical))
(False, False)

df_Numerical

{"type":"dataframe","variable_name":"df_Numerical"}

df_Numerical.shape,df_Categorical.shape
((691406, 78), (692703, 1))

index_numbers_to_fetch = df_Numerical.index.tolist()

df_Categorical = df_Categorical.loc[index_numbers_to_fetch]
```

```

df_Categorical.shape
(691406, 1)
y=pd.get_dummies(data=df_Categorical[ " Label"], drop_first=True)
y
{ "type": "dataframe", "variable_name": "y" }
df_Numerical
{ "type": "dataframe", "variable_name": "df_Numerical" }

```

Min Max Scaler / Standardization

```

df_Numerical_Scaled_data=pd.DataFrame()
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
df_Numerical_Scaled_data[df_Numerical.columns] =
scaler.fit_transform(df_Numerical[df_Numerical.columns])

df_Numerical_Scaled_data
{ "type": "dataframe", "variable_name": "df_Numerical_Scaled_data" }

X=df_Numerical_Scaled_data

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.33, random_state=42)

X_train.shape,y_train.shape,X_test.shape,y_test.shape
((463242, 78), (463242, 1), (228164, 78), (228164, 1))

from sklearn.linear_model import LogisticRegression

# Instantiate logistic regression model with increased max_iter
logistic_model = LogisticRegression(max_iter=1000)

# Fit the model to the training data
logistic_model.fit(X_train, y_train)

/usr/local/lib/python3.10/dist-packages/sklearn/utils/
validation.py:1143: DataConversionWarning: A column-vector y was
passed when a 1d array was expected. Please change the shape of y to
(n_samples, ), for example using ravel().
    y = column_or_1d(y, warn=True)

LogisticRegression(max_iter=1000)

```

```

y_pred = logistic_model.predict(X_test)

from sklearn.metrics import *

accuracy = accuracy_score(y_test,y_pred)*100
confusion_mat = confusion_matrix(y_test,y_pred)
print("Accuracy is",accuracy)
print("Confusion Matrix")
print(confusion_mat)

Accuracy is 96.94430322049053
Confusion Matrix
[[141509  3673]
 [ 3299 79683]]

```

Fitting Different Models

```

from sklearn.svm import SVC
from sklearn.naive_bayes import BernoulliNB
from sklearn import tree
from sklearn.model_selection import cross_val_score
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import *
from sklearn.tree import *
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import BernoulliNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, VotingClassifier

```

Quick Analysis and Results

```

from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report
from sklearn.ensemble import RandomForestClassifier, VotingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import BernoulliNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split

def split_data(X, y, test_size=0.2, random_state=None):
    X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=test_size, random_state=random_state)

```

```

    return X_train, X_test, y_train, y_test

def train_model(model, X_train, y_train):
    y_train = y_train.values.ravel() if len(y_train.shape) > 1 and
y_train.shape[1] == 1 else y_train
    model.fit(X_train, y_train)

def evaluate_model(model, X, y, dataset_name):
    accuracy = accuracy_score(y, model.predict(X))
    confusion_matrix_result = confusion_matrix(y, model.predict(X))
    classification = classification_report(y, model.predict(X))

    print(f'{dataset_name} Model Evaluation')
    print()
    print("Model Accuracy:\n", accuracy)
    print()
    print("Confusion matrix:\n", confusion_matrix_result)
    print()
    print("Classification report:\n", classification)
    print()

X_train, X_test, y_train, y_test = split_data(X, y, test_size=0.2,
random_state=42)

classifiers = [
    ('Naive Baye Classifier', BernoulliNB()),
    ('Decision Tree Classifier',
DecisionTreeClassifier(criterion='entropy', random_state=0)),
    ('Random Forest Classifier',
RandomForestClassifier(n_estimators=100, random_state=0)),
    ('Logistic Regression', LogisticRegression(n_jobs=-1,
random_state=0))
]

models = [
    ('Naive Baye Classifier', BernoulliNB()),
    ('Logistic Regression', LogisticRegression(n_jobs=-1,
random_state=0)),
    ('Voting Classifier', VotingClassifier(estimators=classifiers,
voting='soft'))
]

# Train and evaluate individual models
for model_name, model in models:
    train_model(model, X_train, y_train)
    evaluate_model(model, X_train, y_train, f'Training - {model_name}')
    evaluate_model(model, X_test, y_test, f'Test - {model_name}')

```

Training - Naive Baye Classifier Model Evaluation

Model Accuracy:

0.8217199036744022

Confusion matrix:

```
[[273087  78713]
 [ 19898 181426]]
```

Classification report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| False | 0.93 | 0.78 | 0.85 | 351800 |
| True | 0.70 | 0.90 | 0.79 | 201324 |
| accuracy | | | 0.82 | 553124 |
| macro avg | 0.81 | 0.84 | 0.82 | 553124 |
| weighted avg | 0.85 | 0.82 | 0.82 | 553124 |

Test - Naive Baye Classifier Model Evaluation

Model Accuracy:

0.8237225380020538

Confusion matrix:

```
[[68336 19547]
 [ 4829 45570]]
```

Classification report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| False | 0.93 | 0.78 | 0.85 | 87883 |
| True | 0.70 | 0.90 | 0.79 | 50399 |
| accuracy | | | 0.82 | 138282 |
| macro avg | 0.82 | 0.84 | 0.82 | 138282 |
| weighted avg | 0.85 | 0.82 | 0.83 | 138282 |

Training - Logistic Regression Model Evaluation

Model Accuracy:

0.9721093281072598

Confusion matrix:

```
[[342962  8838]
 [ 6589 194735]]
```

Classification report:

| | precision | recall | f1-score | support |
|--|-----------|--------|----------|---------|
|--|-----------|--------|----------|---------|

| | | | | |
|--------------|------|------|------|--------|
| False | 0.98 | 0.97 | 0.98 | 351800 |
| True | 0.96 | 0.97 | 0.96 | 201324 |
| accuracy | | | 0.97 | 553124 |
| macro avg | 0.97 | 0.97 | 0.97 | 553124 |
| weighted avg | 0.97 | 0.97 | 0.97 | 553124 |

Test - Logistic Regression Model Evaluation

Model Accuracy:
0.9725850074485471

Confusion matrix:
[[85699 2184]
 [1607 48792]]

Classification report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| False | 0.98 | 0.98 | 0.98 | 87883 |
| True | 0.96 | 0.97 | 0.96 | 50399 |
| accuracy | | | 0.97 | 138282 |
| macro avg | 0.97 | 0.97 | 0.97 | 138282 |
| weighted avg | 0.97 | 0.97 | 0.97 | 138282 |

Training - Voting Classifier Model Evaluation

Model Accuracy:
0.9989676817494811

Confusion matrix:
[[351249 551]
 [20 201304]]

Classification report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| False | 1.00 | 1.00 | 1.00 | 351800 |
| True | 1.00 | 1.00 | 1.00 | 201324 |
| accuracy | | | 1.00 | 553124 |
| macro avg | 1.00 | 1.00 | 1.00 | 553124 |
| weighted avg | 1.00 | 1.00 | 1.00 | 553124 |

Test - Voting Classifier Model Evaluation

```

Model Accuracy:
0.998633227751985

Confusion matrix:
[[87707  176]
 [ 13 50386]]

Classification report:
      precision    recall  f1-score   support

     False       1.00     1.00     1.00     87883
      True       1.00     1.00     1.00     50399

  accuracy          1.00          1.00          1.00     138282
  macro avg       1.00     1.00     1.00     138282
weighted avg       1.00     1.00     1.00     138282

```

Stacking

```

def evaluate_model_extended(model, X_train, y_train, X_test, y_test,
dataset_name):
    train_accuracy = accuracy_score(y_train, model.predict(X_train))
    test_accuracy = accuracy_score(y_test, model.predict(X_test))
    f1_score_val = f1_score(y_test, model.predict(X_test))
    precision_val = precision_score(y_test, model.predict(X_test))
    recall_val = recall_score(y_test, model.predict(X_test))

    # Calculate ROC curve and AUC
    y_train_pred_prob = model.predict_proba(X_train)[:, 1]
    fpr_train, tpr_train, _ = roc_curve(y_train, y_train_pred_prob)
    roc_auc_train = roc_auc_score(y_train, y_train_pred_prob)

    y_test_pred_prob = model.predict_proba(X_test)[:, 1]
    fpr_test, tpr_test, _ = roc_curve(y_test, y_test_pred_prob)
    roc_auc_test = roc_auc_score(y_test, y_test_pred_prob)

    # Plot ROC curve
    plt.figure(figsize=(8, 6))
    plt.plot(fpr_train, tpr_train, label=f'Train ROC Curve (AUC = {roc_auc_train:.2f})')
    plt.plot(fpr_test, tpr_test, label=f'Test ROC Curve (AUC = {roc_auc_test:.2f})')
    plt.plot([0, 1], [0, 1], 'k--')
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title(f'ROC Curve - {dataset_name}')

```

```

plt.legend()
plt.show()

# Calculate confusion matrix
train_conf_matrix = confusion_matrix(y_train,
model.predict(X_train))
test_conf_matrix = confusion_matrix(y_test, model.predict(X_test))

# Plot confusion matrix
fig, axes = plt.subplots(1, 2, figsize=(12, 6))

sns.heatmap(train_conf_matrix, annot=True, fmt='d', cmap='Blues',
cbar=False, ax=axes[0])
axes[0].set_title(f'Training Confusion Matrix - {dataset_name}')
axes[0].set_xlabel('Predicted Labels')
axes[0].set_ylabel('True Labels')

sns.heatmap(test_conf_matrix, annot=True, fmt='d', cmap='Blues',
cbar=False, ax=axes[1])
axes[1].set_title(f'Test Confusion Matrix - {dataset_name}')
axes[1].set_xlabel('Predicted Labels')
axes[1].set_ylabel('True Labels')

plt.tight_layout()
plt.show()

# Print evaluation metrics
print(f'{dataset_name} Model Evaluation')
print()
print("Train Accuracy:", train_accuracy)
print("Test Accuracy:", test_accuracy)
print("F1 Score:", f1_score_val)
print("Precision:", precision_val)
print("Recall:", recall_val)
print("ROC AUC (Train):", roc_auc_train)
print("ROC AUC (Test):", roc_auc_test)
print()

# Implement Stacking
base_model_predictions_train = []
base_model_predictions_test = []


# Generate predictions from base models for training and test data
for model_name, model in classifiers:
    train_model(model, X_train, y_train)
    base_model_predictions_train.append(model.predict(X_train))
    base_model_predictions_test.append(model.predict(X_test))


# Train meta-model (Voting Classifier) on predictions from base models
stacked_X_train = np.column_stack(base_model_predictions_train)
stacked_X_test = np.column_stack(base_model_predictions_test)

```

```
stacked_model = VotingClassifier(estimators=classifiers,
voting='soft')
train_model(stacked_model, stacked_X_train, y_train)

# Evaluate stacked model
evaluate_model_extended(stacked_model, stacked_X_train, y_train,
stacked_X_test, y_test, "Stacking")
```





Stacking Model Evaluation

```
Train Accuracy: 0.9999367230494428
Test Accuracy: 0.9997035044329703
F1 Score: 0.9995932902816218
Precision: 0.9994842293195795
Recall: 0.999702375047124
ROC AUC (Train): 0.9999619980058335
ROC AUC (Test): 0.9998862821145618
```

```
for model_name, model in models:
    if 'Decision Tree' in model_name or 'Random Forest' in model_name:
        importances = model.feature_importances_
        feature_names = X.columns
        indices = np.argsort(importances)[::-1]

        plt.figure(figsize=(10, 6))
        plt.title(f'Feature Importance for {model_name}')
        plt.bar(range(X.shape[1]), importances[indices],
align='center')
        plt.xticks(range(X.shape[1]), [feature_names[i] for i in
indices], rotation=90)
        plt.xlim([-1, X.shape[1]])
        plt.show()
    elif 'Voting Classifier' in model_name:
        # Voting classifier doesn't have feature importances
        print(f"No feature importances for {model_name}")
    else:
        # Handle other models if needed
        pass
```

```
No feature importances for Voting Classifier
```

```
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import *
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import BernoulliNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import RandomForestClassifier, VotingClassifier
```

Visualisations Using ROC & Precision-Recall Curve

```
def split_data(X, y, test_size=0.2, random_state=None):
    X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=test_size, random_state=random_state)
    return X_train, X_test, y_train, y_test

def train_model(model, X_train, y_train):
    model.fit(X_train, y_train)

def evaluate_model(model, X, y):
    y = np.ravel(y)

    X_np = np.array(X)

    y_pred = model.predict(X_np)

    accuracy = accuracy_score(y, y_pred)
    confusion_matrix_result = confusion_matrix(y, y_pred)
    classification = classification_report(y, y_pred)
    f1 = f1_score(y, y_pred)
    precision = precision_score(y, y_pred)
    recall = recall_score(y, y_pred)

    return {
        'accuracy': accuracy,
        'confusion_matrix': confusion_matrix_result,
        'classification_report': classification,
        'f1_score': f1,
        'precision': precision,
        'recall': recall
    }
```

```

def plot_roc_curve(fpr, tpr, auc, model_name):
    plt.figure(figsize=(8, 6))
    plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve
(area = {:.2f})'.format(auc))
    plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('Receiver Operating Characteristic for
{}'.format(model_name))
    plt.legend(loc='lower right')
    plt.show()

def plot_precision_recall_curve(precision, recall, auc, model_name):
    plt.figure(figsize=(8, 6))
    plt.plot(recall, precision, color='darkorange', lw=2, label='PR
curve (area = {:.2f})'.format(auc))
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('Recall')
    plt.ylabel('Precision')
    plt.title('Precision-Recall Curve for {}'.format(model_name))
    plt.legend(loc='lower left')
    plt.show()

models = [
    ('Naive Baye Classifier', BernoulliNB()),
    ('Random Forest Classifier',
RandomForestClassifier(n_estimators=100, random_state=0)),
    ('Logistic Regression', LogisticRegression(n_jobs=-1,
random_state=0)),
    ('Voting Classifier', VotingClassifier(estimators=classifiers,
voting='soft'))
]

results_df = pd.DataFrame(columns=['Model', 'Train Accuracy', 'Test
Accuracy', 'F1 Score', 'Precision', 'Recall'])

for model_name, model in models:
    train_model(model, X_train, y_train)
    train_evaluation = evaluate_model(model, X_train, y_train)
    test_evaluation = evaluate_model(model, X_test, y_test)

    results_df = pd.concat([results_df, pd.DataFrame({
        'Model': [model_name],
        'Train Accuracy': [train_evaluation['accuracy']],
        'Test Accuracy': [test_evaluation['accuracy']],
        'F1 Score': [test_evaluation['f1_score']]})])

```

```
'Precision': [test_evaluation['precision']],
'Recall': [test_evaluation['recall']],
})], ignore_index=True)

if model_name != 'Voting Classifier':
    y_prob = model.predict_proba(X_test)[:, 1]
    fpr, tpr, _ = roc_curve(y_test, y_prob)
    roc_auc = auc(fpr, tpr)


    precision, recall, _ = precision_recall_curve(y_test, y_prob)
    pr_auc = auc(recall, precision)


    plot_roc_curve(fpr, tpr, roc_auc, model_name)
    plot_precision_recall_curve(precision, recall, pr_auc,
model_name)

print(results_df)

/usr/local/lib/python3.10/dist-packages/sklearn/utils/
validation.py:1143: DataConversionWarning: A column-vector y was
passed when a 1d array was expected. Please change the shape of y to
(n_samples, ), for example using ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439:
UserWarning: X does not have valid feature names, but BernoulliNB was
fitted with feature names
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439:
UserWarning: X does not have valid feature names, but BernoulliNB was
fitted with feature names
    warnings.warn(
```


Receiver Operating Characteristic for Naive Baye Classifier






```
<ipython-input-48-c8a1ed5953ad>:6: DataConversionWarning: A column-  
vector y was passed when a 1d array was expected. Please change the  
shape of y to (n_samples,), for example using ravel().  
    model.fit(X_train, y_train)  
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439:  
UserWarning: X does not have valid feature names, but  
RandomForestClassifier was fitted with feature names  
    warnings.warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439:  
UserWarning: X does not have valid feature names, but  
RandomForestClassifier was fitted with feature names  
    warnings.warn(
```


Receiver Operating Characteristic for Random Forest Classifier






```
/usr/local/lib/python3.10/dist-packages/sklearn/utils/
validation.py:1143: DataConversionWarning: A column-vector y was
passed when a 1d array was expected. Please change the shape of y to
(n_samples, ), for example using ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439:
UserWarning: X does not have valid feature names, but
LogisticRegression was fitted with feature names
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439:
UserWarning: X does not have valid feature names, but
LogisticRegression was fitted with feature names
    warnings.warn(
```

Receiver Operating Characteristic for Logistic Regression





```
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.py:99: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:134: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439:
UserWarning: X does not have valid feature names, but BernoulliNB was
fitted with feature names
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439:
UserWarning: X does not have valid feature names, but
DecisionTreeClassifier was fitted with feature names
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439:
UserWarning: X does not have valid feature names, but
```

```
RandomForestClassifier was fitted with feature names
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439:
UserWarning: X does not have valid feature names, but
LogisticRegression was fitted with feature names
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439:
UserWarning: X does not have valid feature names, but BernoulliNB was
fitted with feature names
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439:
UserWarning: X does not have valid feature names, but
DecisionTreeClassifier was fitted with feature names
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439:
UserWarning: X does not have valid feature names, but
RandomForestClassifier was fitted with feature names
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439:
UserWarning: X does not have valid feature names, but
LogisticRegression was fitted with feature names
    warnings.warn(
```

| Score \ Model | Train Accuracy | Test Accuracy | F1 |
|----------------------------|----------------|---------------|----------|
| 0 Naive Baye Classifier | 0.821720 | 0.823723 | 0.788982 |
| 1 Random Forest Classifier | 0.999669 | 0.999573 | 0.999415 |
| 2 Logistic Regression | 0.972109 | 0.972585 | 0.962604 |
| 3 Voting Classifier | 0.998968 | 0.998633 | 0.998128 |

| | Precision | Recall |
|---|-----------|----------|
| 0 | 0.699817 | 0.904185 |
| 1 | 0.999108 | 0.999722 |
| 2 | 0.957156 | 0.968114 |
| 3 | 0.996519 | 0.999742 |

```
accuracy_metrics = ['Train Accuracy', 'Test Accuracy']
model_names = results_df['Model']
num_models = len(models)

bar_width = 0.35

indices = np.arange(len(model_names))

colors = ['b', 'g']

plt.figure(figsize=(12, 6))
```

```

for i, metric in enumerate(accuracy_metrics):
    x_positions = [ind + i * bar_width for ind in indices]


    metric_values = results_df[metric]

    plt.bar(x_positions, metric_values, width=bar_width, label=metric,
color=colors[i])

plt.xlabel('Model')
plt.ylabel('Accuracy')
plt.title('Training and Testing Accuracy Comparison')
plt.xticks(indices + (bar_width * (num_models - 1)) / 2, model_names,
rotation=45)
plt.legend(loc='lower right')

plt.tight_layout()
plt.show()

```




```


for model_name, model in models:
    y_pred = model.predict(X_test)
    cm = confusion_matrix(y_test, y_pred)
    disp = ConfusionMatrixDisplay(confusion_matrix=cm,
display_labels=["Negative", "Positive"])
    disp.plot(cmap=plt.cm.Blues, values_format='d')
    plt.title(f'Confusion Matrix for {model_name}')
    plt.show()

```


Confusion Matrix for Naive Baye Classifier




Confusion Matrix for Random Forest Classifier



Confusion Matrix for Logistic Regression






```

for model_name, model in models:
    if 'Decision Tree' in model_name or 'Random Forest' in model_name:
        importances = model.feature_importances_
        feature_names = X.columns
        indices = np.argsort(importances)[::-1]

        plt.figure(figsize=(10, 6))
        plt.title(f'Feature Importance for {model_name}')
        plt.bar(range(X.shape[1]), importances[indices],
align='center')
        plt.xticks(range(X.shape[1]), [feature_names[i] for i in
indices], rotation=90)
        plt.xlim([-1, X.shape[1]])
        plt.show
    
```



Visualize learning curves for each model

```
from sklearn.model_selection import learning_curve

for model_name, model in models:
    train_sizes, train_scores, test_scores = learning_curve(model,
X_train, y_train, cv=5, scoring='accuracy')
    train_scores_mean = np.mean(train_scores, axis=1)
    test_scores_mean = np.mean(test_scores, axis=1)

    plt.figure(figsize=(8, 6))
    plt.title(f'Learning Curve for {model_name}')
    plt.xlabel("Training examples")
    plt.ylabel("Score")
    plt.grid()
    plt.plot(train_sizes, train_scores_mean, 'o-', label="Training
```


```
score")
    plt.plot(train_sizes, test_scores_mean, 'o-', label="Cross-validation score")
    plt.legend(loc="best")
    plt.show()

/usr/local/lib/python3.10/dist-packages/sklearn/utils/
validation.py:1143: DataConversionWarning: A column-vector y was
passed when a 1d array was expected. Please change the shape of y to
(n_samples, ), for example using ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
```

```
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
```

```
example using ravel().  
    y = column_or_1d(y, warn=True)  
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11  
43: DataConversionWarning: A column-vector y was passed when a 1d  
array was expected. Please change the shape of y to (n_samples, ), for  
example using ravel().  
    y = column_or_1d(y, warn=True)  
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11  
43: DataConversionWarning: A column-vector y was passed when a 1d  
array was expected. Please change the shape of y to (n_samples, ), for  
example using ravel().  
    y = column_or_1d(y, warn=True)  
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11  
43: DataConversionWarning: A column-vector y was passed when a 1d  
array was expected. Please change the shape of y to (n_samples, ), for  
example using ravel().  
    y = column_or_1d(y, warn=True)  
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11  
43: DataConversionWarning: A column-vector y was passed when a 1d  
array was expected. Please change the shape of y to (n_samples, ), for  
example using ravel().  
    y = column_or_1d(y, warn=True)  
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11  
43: DataConversionWarning: A column-vector y was passed when a 1d  
array was expected. Please change the shape of y to (n_samples, ), for  
example using ravel().  
    y = column_or_1d(y, warn=True)
```

Learning Curve for Naive Baye Classifier




```
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:686: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:686: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:686: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:686: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:686: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
```

```
ation.py:686: DataConversionWarning: A column-vector y was passed when
a 1d array was expected. Please change the shape of y to (n_samples,),
for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_valid
ation.py:686: DataConversionWarning: A column-vector y was passed when
a 1d array was expected. Please change the shape of y to (n_samples,),
for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_valid
ation.py:686: DataConversionWarning: A column-vector y was passed when
a 1d array was expected. Please change the shape of y to (n_samples,),
for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_valid
ation.py:686: DataConversionWarning: A column-vector y was passed when
a 1d array was expected. Please change the shape of y to (n_samples,),
for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_valid
ation.py:686: DataConversionWarning: A column-vector y was passed when
a 1d array was expected. Please change the shape of y to (n_samples,),
for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_valid
ation.py:686: DataConversionWarning: A column-vector y was passed when
a 1d array was expected. Please change the shape of y to (n_samples,),
for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_valid
ation.py:686: DataConversionWarning: A column-vector y was passed when
a 1d array was expected. Please change the shape of y to (n_samples,),
for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_valid
ation.py:686: DataConversionWarning: A column-vector y was passed when
a 1d array was expected. Please change the shape of y to (n_samples,),
for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_valid
ation.py:686: DataConversionWarning: A column-vector y was passed when
a 1d array was expected. Please change the shape of y to (n_samples,),
for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:686: DataConversionWarning: A column-vector y was passed when
a 1d array was expected. Please change the shape of y to (n_samples,), 
for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:686: DataConversionWarning: A column-vector y was passed when
a 1d array was expected. Please change the shape of y to (n_samples,), 
for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:686: DataConversionWarning: A column-vector y was passed when
a 1d array was expected. Please change the shape of y to (n_samples,), 
for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:686: DataConversionWarning: A column-vector y was passed when
a 1d array was expected. Please change the shape of y to (n_samples,), 
for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:686: DataConversionWarning: A column-vector y was passed when
a 1d array was expected. Please change the shape of y to (n_samples,), 
for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:686: DataConversionWarning: A column-vector y was passed when
a 1d array was expected. Please change the shape of y to (n_samples,), 
for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:686: DataConversionWarning: A column-vector y was passed when
a 1d array was expected. Please change the shape of y to (n_samples,), 
for example using ravel().
    estimator.fit(X_train, y_train, **fit_params)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:686: DataConversionWarning: A column-vector y was passed when
a 1d array was expected. Please change the shape of y to (n_samples,), 
for example using ravel().
```

```
estimator.fit(X_train, y_train, **fit_params)
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py:686: DataConversionWarning: A column-vector y was passed when
a 1d array was expected. Please change the shape of y to (n_samples,),
for example using ravel().
estimator.fit(X_train, y_train, **fit_params)
```




```
/usr/local/lib/python3.10/dist-packages/sklearn/utils/
validation.py:1143: DataConversionWarning: A column-vector y was
passed when a 1d array was expected. Please change the shape of y to
(n_samples, ), for example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
```

```
example using ravel().  
    y = column_or_1d(y, warn=True)  
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11  
43: DataConversionWarning: A column-vector y was passed when a 1d  
array was expected. Please change the shape of y to (n_samples, ), for  
example using ravel().  
    y = column_or_1d(y, warn=True)  
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11  
43: DataConversionWarning: A column-vector y was passed when a 1d  
array was expected. Please change the shape of y to (n_samples, ), for  
example using ravel().  
    y = column_or_1d(y, warn=True)  
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11  
43: DataConversionWarning: A column-vector y was passed when a 1d  
array was expected. Please change the shape of y to (n_samples, ), for  
example using ravel().  
    y = column_or_1d(y, warn=True)  
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11  
43: DataConversionWarning: A column-vector y was passed when a 1d  
array was expected. Please change the shape of y to (n_samples, ), for  
example using ravel().  
    y = column_or_1d(y, warn=True)  
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11  
43: DataConversionWarning: A column-vector y was passed when a 1d  
array was expected. Please change the shape of y to (n_samples, ), for  
example using ravel().  
    y = column_or_1d(y, warn=True)  
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11  
43: DataConversionWarning: A column-vector y was passed when a 1d  
array was expected. Please change the shape of y to (n_samples, ), for  
example using ravel().  
    y = column_or_1d(y, warn=True)  
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11  
43: DataConversionWarning: A column-vector y was passed when a 1d  
array was expected. Please change the shape of y to (n_samples, ), for  
example using ravel().  
    y = column_or_1d(y, warn=True)  
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11  
43: DataConversionWarning: A column-vector y was passed when a 1d  
array was expected. Please change the shape of y to (n_samples, ), for  
example using ravel().
```

```
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
```

```
43: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
y = column_or_1d(y, warn=True)
```



```
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.py:99: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
y = column_or_1d(y, warn=True)
```


```
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:134: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:99: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:134: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:99: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:134: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:99: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:134: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:99: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, warn=True)
```

```
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:134: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:99: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:134: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:99: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:134: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:99: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:134: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:99: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:134: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:99: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
```

```
example using ravel().  
    y = column_or_1d(y, warn=True)  
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p  
y:134: DataConversionWarning: A column-vector y was passed when a 1d  
array was expected. Please change the shape of y to (n_samples, ), for  
example using ravel().  
    y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)  
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p  
y:99: DataConversionWarning: A column-vector y was passed when a 1d  
array was expected. Please change the shape of y to (n_samples, ), for  
example using ravel().  
    y = column_or_1d(y, warn=True)  
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p  
y:134: DataConversionWarning: A column-vector y was passed when a 1d  
array was expected. Please change the shape of y to (n_samples, ), for  
example using ravel().  
    y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)  
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p  
y:99: DataConversionWarning: A column-vector y was passed when a 1d  
array was expected. Please change the shape of y to (n_samples, ), for  
example using ravel().  
    y = column_or_1d(y, warn=True)  
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p  
y:134: DataConversionWarning: A column-vector y was passed when a 1d  
array was expected. Please change the shape of y to (n_samples, ), for  
example using ravel().  
    y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)  
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p  
y:99: DataConversionWarning: A column-vector y was passed when a 1d  
array was expected. Please change the shape of y to (n_samples, ), for  
example using ravel().  
    y = column_or_1d(y, warn=True)  
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p  
y:134: DataConversionWarning: A column-vector y was passed when a 1d  
array was expected. Please change the shape of y to (n_samples, ), for  
example using ravel().  
    y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)  
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p  
y:99: DataConversionWarning: A column-vector y was passed when a 1d
```

```
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:134: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:99: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:134: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:99: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:134: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:99: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:134: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
```

```
y:99: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:134: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:99: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:134: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:99: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:134: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:99: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.p
y:134: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)
```



Create a grouped bar plot for comparing multiple metrics across models

```
metrics_to_compare = ['Train Accuracy', 'Test Accuracy', 'F1 Score',
'Precision', 'Recall']
model_names = results_df['Model']
num_models = len(models)

bar_width = 0.15

indices = np.arange(len(model_names))

colors = ['b', 'g', 'm', 'y', 'c']

plt.figure(figsize=(12, 6))

for i, metric in enumerate(metrics_to_compare):
    x_positions = [ind + i * bar_width for ind in indices]
```

```

metric_values = results_df[metric]

plt.bar(x_positions, metric_values, width=bar_width, label=metric,
color=colors[i])

plt.xlabel('Model')
plt.ylabel('Score')
plt.title('Model Comparison')
plt.xticks(indices + (bar_width * (num_models - 1)) / 2, model_names,
rotation=45)
plt.legend(loc='lower right')

plt.tight_layout()
plt.show()

```

